



Mutex Locks

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}
```

```
release() {  
    available = true;  
}
```

Problem: *spinlock* - process "spins" while waiting for the lock.





Semaphore

- Semaphore S – integer variable
- Two standard operations modify S : `wait()` and `signal()`
 - Originally called $P() \rightarrow P$ (from the Dutch *proberen*, "to test");
 - $V() \rightarrow V$ (from *verhogen*, "to increment").
- Less complicated
- Can only be accessed via two indivisible (atomic) operations
 - `wait (S) {`
 - `while S <= 0`
 - `; // no-op`
 - `S--;`
 - `}`
 - `signal (S) {`
 - `S++;`
 - `}`

2 versions of semaphore exists, 1st version with busy waiting, 2nd version with blocking call





Semaphore as General Synchronization Tool

- **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement
 - Also known as **mutex locks**
- Mutual-exclusion implementation with **Binary** semaphore
- Semaphore mutex; // initialized to 1

```
do {  
    wait (mutex);  
    // Critical Section  
    signal (mutex);  
    // remainder section  
} while (TRUE);
```





- **Use a binary semaphore** when the goal is to protect a critical section. The binary semaphore ensures that only one thread or process can execute in the critical section at any given time.
- **Use a counting semaphore** when you have a fixed number of resources that can be shared among multiple threads or processes, but the access to these resources does not constitute a critical section that requires mutual exclusion.
 - **Managing access to multiple resources.** Counting semaphores allow multiple threads or processes to access shared resources concurrently, up to a specified limit. They are not typically used to enforce mutual exclusion in critical sections, as they allow more than one thread to proceed if the count is greater than 1.
 - Eg: DB connections





- **Counting** semaphore – integer value can range over an unrestricted domain
- **Case 1: To control access to a given resource**
- **Semaphore is initialized to the number of resources**
- Semaphore Count = 0; all resources are utilized
- **Case 2 : process synchronization**
- (if we want the restriction on the order of execution: s2 is executed only after s1 completes)
- **synch=0**

```
S1;  
signal(synch);
```

```
wait(synch);  
S2;
```

Disadvantage: Busy waiting thus wastes CPU cycles





- While a process is in its critical section, any other process that tries to enter its critical section must loop continuously in the entry code.
- This type of semaphore is also called **spinlock** because the process "spins" while waiting for the lock.
- Busy waiting wastes CPU cycles that some other process might be able to use productively.
- To overcome the need for busy waiting, we can modify the definition of the wait() and signal() semaphore operations.
- rather than engaging in busy waiting, the **process can block itself**. The block operation places a process into a waiting queue associated with the semaphore





Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```





Semaphore Implementation with no Busy waiting (Cont.)

- Synchronization tool that does not require busy waiting

- Implementation of wait:

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

Entry

Here, semaphore values may be negative, whereas semaphore values are never negative under the classical definition of semaphores with busy waiting.

Semaphore value -ve indicates the processes waiting to acquire that semaphore

- Implementation of signal:

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

Put the processes from the blocked list to ready queue... not into CS

Exit





Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event
- Let **S** and **Q** be two semaphores initialized to 1

P_0
wait (S); **S=0**
wait (Q); **Q=-1**
.
.
.
signal (S);
signal (Q);

P_1
wait (Q); **Q=0**
wait (S); **S=-1**
.
.
.
signal (Q);
signal (S);

- **Starvation** – indefinite blocking
 - occur if we remove processes from the list associated with a semaphore in LIFO order





Priority Inversion

- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
- Solved via **priority-inheritance protocol**
- L and H share/modify kernel data(CS)
- Assume we have three processes, L , M , and H , whose priorities follow the order $L < M < H$. Assume that process H requires resource R , which is currently being accessed by process L . Ordinarily, process H would wait for L to finish using resource R .
- M becomes runnable, thereby preempting process
- a process with a lower priority—process M —has affected how long process H must wait for L to relinquish resource R .
- This problem is known as **priority inversion**.
- priority-inheritance protocol would allow process L to temporarily inherit the priority of process H , thereby preventing process M from preempting its execution. When process L had finished using resource R , it would relinquish its inherited priority from H and assume its original priority. Because resource R would now be available, process H —not M —would run next.

**L and H(not M)
share resource R**

